

Introduction to Basic and Advanced Statistical Modelling

From LM to GLM

Alexandre Courtiol

Table of contents

Linear Models (LM)

Interpreting coefficients

Generalised Linear Models (GLM)

Tests

Assumptions

Model construction

A simple habitat occupancy model

Linear Models (LM)

LM — what is a statistical model?

Statistical model

an approximate description of the *data generating process*, i.e. an approximation of how the world generates data

LM — what is a statistical model?

Statistical model

an approximate description of the *data generating process*, i.e. an approximation of how the world generates data

NB: seeing models as a *data generating process* may not be intuitive at first, but it is both conceptually and practically crucial!

In particular, it helps

- correctly interpreting model parameters
- planning sample size before data collection
- testing the coverage of confidence intervals (i.e. do they work as intended)
- testing the robustness of the statistical test (i.e. check the rate of false positives)
- performing prior and posterior predictive checks during Bayesian analyses

LM — what is a linear model (LM)?

Linear model

a data generating process produced by a *linear function*: the data are generated from a set of terms by multiplying each term by a constant (a model parameter) and summing them

LM — what is a linear model (LM)?

Linear model

a data generating process produced by a *linear function*: the data are generated from a set of terms by multiplying each term by a constant (a model parameter) and summing them

NB:

- linear models encompass the classic linear regression which forms a straight line, but also many other models (e.g. polynomials are linear and yet do not form straight lines) → LM are linear in parameters (not necessarily in variables)
- much of the difficulty people encounter when dealing with more complex models (e.g. Generalised Linear Models, Linear Mixed-effects Models, Hierarchical Models) is often caused by not having sufficiently understood the LM

LM — notations of a LM: simple notation

$$y_i = \beta_0 + \beta_1 \times x_{1,i} + \beta_2 \times x_{2,i} + \cdots + \beta_p \times x_{p,i} + \epsilon_i \text{ with } \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

LM — notations of a LM: simple notation

$$y_i = \beta_0 + \beta_1 \times x_{1,i} + \beta_2 \times x_{2,i} + \dots + \beta_p \times x_{p,i} + \epsilon_i \text{ with } \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- y_i = the response variable / dependent variable / observations to explain
- β_j = the model parameters / regression coefficients
- $x_{j,i}$ = regressors / constants derived from the predictor variables (aka predictors for short)
- ϵ_i = the errors / residual errors

LM — notations of a LM: simple notation

$$y_i = \beta_0 + \beta_1 \times x_{1,i} + \beta_2 \times x_{2,i} + \dots + \beta_p \times x_{p,i} + \epsilon_i \text{ with } \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- y_i = the response variable / dependent variable / observations to explain
- β_j = the model parameters / regression coefficients
- $x_{j,i}$ = regressors / constants derived from the predictor variables (aka predictors for short)
- ϵ_i = the errors / residual errors

NB:

- this is the notation that biologists should use in their paper/thesis to describe their model
- the regressors are sometimes directly given by the columns from your dataset (i.e. the predictors), but not always (more on that later)
- when parameters are not known but estimated they should wear a hat (e.g. $\hat{\beta}_0$) and ϵ_i (also sometimes called ε_i) usually becomes e_i or r_i

LM — notations of a LM: usual R notation

$$y \sim x_1 + x_2 + \dots + x_p$$

NB:

- the parameters and the error term are not written down

LM — notations of a LM: matrix notation

$$Y = X\beta + \epsilon$$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,p} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,p} \\ 1 & x_{3,1} & x_{3,2} & \dots & x_{3,p} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_{n,1} & x_{n,2} & \dots & x_{n,p} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \dots \\ \beta_p \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \dots \\ \epsilon_n \end{bmatrix}$$

LM — notations of a LM: matrix notation

$$Y = X\beta + \epsilon$$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,p} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,p} \\ 1 & x_{3,1} & x_{3,2} & \dots & x_{3,p} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_{n,1} & x_{n,2} & \dots & x_{n,p} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \dots \\ \beta_p \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \dots \\ \epsilon_n \end{bmatrix}$$

- Y = the vector of observations
- X = a matrix called the design (or model) matrix
- β = the vector of model parameters
- ϵ = the vector of errors, with again $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$

NB: this is the notation for statisticians and computer scientists, which is particularly useful when interpreting coefficients

LM — data generating process: data simulation

Let's enter **God mode** → we pretend that **Emerald Lion** (a goddess) knows how her world generates data; she informed us that in her world aliens keep growing with the number of humans they eat as follows:

$$\begin{cases} \text{size}_i = 50 + 1.5 \times \text{humans_eaten}_i + \epsilon_i \\ \epsilon_i \sim \mathcal{N}(0, \sigma^2 = 4) \end{cases}$$

LM — data generating process: data simulation

Let's enter **God mode** → we pretend that **Emerald Lion** (a goddess) knows how her world generates data; she informed us that in her world aliens keep growing with the number of humans they eat as follows:

$$\begin{cases} \text{size}_i = 50 + 1.5 \times \text{humans_eaten}_i + \epsilon_i \\ \epsilon_i \sim \mathcal{N}(0, \sigma^2 = 4) \end{cases}$$

We can thus generate the size of 12 aliens that have eaten 1–3 humans following E.L.'s recipe:

```
set.seed(2)
N <- 12
Alien <- tibble(humans_eaten = rep(1:3, each = 4), # alternative e.g. rpois() would do
                intercept = 50, slope = 1.5,
                error = rnorm(n = N, mean = 0, sd = 2), # must be rnorm!
                size = intercept + slope*humans_eaten + error)
```

LM — data generating process: data simulation

Let's enter **God mode** → we pretend that **Emerald Lion** (a goddess) knows how her world generates data; she informed us that in her world aliens keep growing with the number of humans they eat as follows:

$$\begin{cases} \text{size}_i = 50 + 1.5 \times \text{humans_eaten}_i + \epsilon_i \\ \epsilon_i \sim \mathcal{N}(0, \sigma^2 = 4) \end{cases}$$

We can thus generate the size of 12 aliens that have eaten 1–3 humans following E.L.'s recipe:

```
set.seed(2)
N <- 12
Alien <- tibble(humans_eaten = rep(1:3, each = 4), # alternative e.g. rpois() would do
                intercept = 50, slope = 1.5,
                error = rnorm(n = N, mean = 0, sd = 2), # must be rnorm!
                size = intercept + slope*humans_eaten + error)
```

NB: we could also write the last 2 lines as:

```
size <- rnorm(n = 12, mean = intercept + slope*humans_eaten, sd = 2)
```


LM — data generating process: data check

Simulated data

Alien

```
# A tibble: 12 x 5
  humans_eaten intercept slope  error  size
      <int>      <dbl> <dbl>  <dbl> <dbl>
1           1         50   1.5 -1.79  49.7
2           1         50   1.5  0.370  51.9
3           1         50   1.5  3.18  54.7
4           1         50   1.5 -2.26  49.2
5           2         50   1.5 -0.161  52.8
# i 7 more rows
```

- which columns show the *model parameters*?
- which columns show *regressor* (here = predictor)?
- which columns show *errors*?
- which columns show *response variable*?

Practice 1

Simulated data

Alien

```
# A tibble: 12 x 5
  humans_eaten intercept slope  error  size
      <int>      <dbl> <dbl>  <dbl> <dbl>
1           1         50   1.5 -1.79  49.7
2           1         50   1.5  0.370 51.9
3           1         50   1.5  3.18  54.7
4           1         50   1.5 -2.26  49.2
5           2         50   1.5 -0.161 52.8
# i 7 more rows
```

- which columns show the *model parameters*?
- which columns show *regressor* (here = predictor)?
- which columns show *errors*?
- which columns show *response variable*?
- what would the *mean response value* be, for an alien who does not eat humans?
- what would the *mean response value* be, for an alien who eats 2 humans?

Solution

- what would the *mean response value* be, for an alien who does not eat humans?

$$\text{mean}(y \mid \text{humans_eaten} = 0) = 50 + 1.5 \times 0 = 50$$

Solution

- what would the *mean response value* be, for an alien who does not eat humans?

$$\text{mean}(y \mid \text{humans_eaten} = 0) = 50 + 1.5 \times 0 = 50$$

- what would the *mean response value* be, for an alien who eats 2 humans?

$$\text{mean}(y \mid \text{humans_eaten} = 2) = 50 + 1.5 \times 2 = 53$$

LM — model fitting: frequentist approach

In practice, model parameters are **unknown**, so we have to estimate them

→ that is the whole point of statistical modelling... → the God mode is now off

LM — model fitting: frequentist approach

In practice, model parameters are *unknown*, so we have to estimate them

→ that is the whole point of statistical modelling... → the God mode is now off

- in frequentist statistics, *fitting the model* means estimating the values of the model parameters that maximise the probability of the data

```
fit_alien_lm <- lm(size ~ humans_eaten, data = Alien)
summary(fit_alien_lm) # or coef(fit_alien_lm) for a focus on coefficients only
```

Call:

```
lm(formula = size ~ humans_eaten, data = Alien)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.9710	-1.1744	-0.4725	0.7023	3.4655

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	48.8353	1.3454	36.298	5.99e-12 ***
humans_eaten	2.3749	0.6228	3.813	0.00341 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.762 on 10 degrees of freedom

Multiple R-squared: 0.5925, Adjusted R-squared: 0.5518

F-statistic: 14.54 on 1 and 10 DF, p-value: 0.003411

LM — model fitting: interlude on ordinary least squares (OLS)

In LM (but not GLM), finding the estimates that maximise the probability of the data given the model (i.e. maximising the likelihood) is equivalent to minimising the sum of the squared residuals, so this is what `lm()` actually does since it is computationally simpler than computing the likelihood

To best illustrate the OLS concept, we need to simulate data a little more spread out


```
set.seed(123)
Alien2 <- tibble(humans_eaten = 1:12, intercept = 50, slope = 1.5,
                 size = rnorm(n = 12, mean = intercept + slope*humans_eaten, sd = 4))
fit_alien2_lm <- lm(size ~ humans_eaten, data = Alien2)
```

LM — model fitting: interlude on ordinary least squares (OLS)

In LM (but not GLM), finding the estimates that maximise the probability of the data given the model (i.e. maximising the likelihood) is equivalent to minimising the sum of the squared residuals, so this is what `lm()` actually does since it is computationally simpler than computing the likelihood

To best illustrate the OLS concept, we need to simulate data a little more spread out

```
set.seed(123)
Alien2 <- tibble(humans_eaten = 1:12, intercept = 50, slope = 1.5,
                 size = rnorm(n = 12, mean = intercept + slope*humans_eaten, sd = 4))
fit_alien2_lm <- lm(size ~ humans_eaten, data = Alien2)
```




Residuals on best fit

LM — model fitting: interlude on ordinary least squares (OLS)


In LM (but not GLM), finding the estimates that maximise the probability of the data given the model (i.e. maximising the likelihood) is equivalent to minimising the sum of the squared residuals, so this is what `lm()` actually does since it is computationally simpler than computing the likelihood

To best illustrate the OLS concept, we need to simulate data a little more spread out

```
set.seed(123)
Alien2 <- tibble(humans_eaten = 1:12, intercept = 50, slope = 1.5,
                 size = rnorm(n = 12, mean = intercept + slope*humans_eaten, sd = 4))
fit_alien2_lm <- lm(size ~ humans_eaten, data = Alien2)
```



Residuals on best fit




Squared residuals on best fit

LM — model fitting: interlude on ordinary least squares (OLS)


In LM (but not GLM), finding the estimates that maximise the probability of the data given the model (i.e. maximising the likelihood) is equivalent to minimising the sum of the squared residuals, so this is what `lm()` actually does since it is computationally simpler than computing the likelihood

To best illustrate the OLS concept, we need to simulate data a little more spread out

```
set.seed(123)
Alien2 <- tibble(humans_eaten = 1:12, intercept = 50, slope = 1.5,
                 size = rnorm(n = 12, mean = intercept + slope*humans_eaten, sd = 4))
fit_alien2_lm <- lm(size ~ humans_eaten, data = Alien2)
```



Squared residuals on bad fit




Squared residuals on best fit

LM — model fitting: fit vs ground truth, & vocabulary

```
Alien |> # we are back to using our original data & fit!
```

```
mutate(mean_response_values = intercept + slope * humans_eaten, # truth in population  
       fitted_values = fitted(fit_alien_lm),  
       residuals = residuals(fit_alien_lm)) -> Alien
```




- a *prediction* is the average value of the response given by the fit for a large number of observations with the same predictor characteristics (i.e. any point on the coral line is a prediction)
- a *fitted value* is a *prediction* at observed predictor characteristics (i.e. only the points for 1, 2 and 3 humans on the coral line)
- a *residual* is a difference between an observation and its fitted value (i.e. $observation_i = fitted_i + residual_i$)

LM — model fitting: Bayesian approach with {brms}

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Let's first use {brms}, which is simpler than {nimble} but handles a smaller set of models

```
library(brms) # Bayesian regression models using stan
fit_alien_lm_brms <- brm(size ~ humans_eaten, data = Alien)
mcmc_plot(fit_alien_lm_brms, type = "trace") # many diagnostic plots can be produced
```




LM — model fitting: Bayesian approach with {brms}

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Here are the posteriors distributions

```
mcmc_plot(fit_alien_lm_brms, type = "dens_overlay") # many diagnostic plots can be produced
```



LM — model fitting: Bayesian approach with {brms}

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Here is the summary for the posteriors distributions

```
summary(fit_alien_lm_brms) # or fixef(fit_alien_lm_brms) for a focus on coefficients
```

```
Family: gaussian
Links: mu = identity
Formula: size ~ humans_eaten
Data: Alien (Number of observations: 12)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000
```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	48.83	1.51	45.81	51.85	1.00	3306	2454
humans_eaten	2.38	0.70	0.99	3.76	1.00	3166	2281

Further Distributional Parameters:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	1.92	0.46	1.26	3.03	1.00	2388	2555

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

LM — model fitting: Bayesian approach with {nimble}

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Let's now use {nimble}

```
library(nimble)
library(posterior)
library(bayesplot)

code <- nimbleCode({
  for (i in seq_len(N)) { # model mimicking the data generating process
    size[i] ~ dnorm(mean = intercept + slope*humans_eaten[i], sd = sigma)
  }
  # priors
  intercept ~ dnorm(mean = 50, sd = 5)
  slope ~ dnorm(mean = 0, sd = 5)
  sigma ~ dunif(min = 0, max = 5)
})

constants <- list(N = nrow(Alien), humans_eaten = Alien$humans_eaten)

inits <- list(intercept = mean(Alien$size), slope = 0, sigma = 1)
```

LM — model fitting: prior predictive checking

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Before fitting the model, let's make sure our priors are not too vague

```
set.seed(123)
n_sims <- 1000

prior <- tibble(sim = seq_len(n_sims),
               intercept = rnorm(n_sims, mean = 50, sd = 5),
               slope = rnorm(n_sims, mean = 0, sd = 5),
               sigma = runif(n_sims, min = 0, max = 5))

humans_eaten_possible <- min(Alien$humans_eaten):max(Alien$humans_eaten)

expand_grid(prior, humans_eaten = humans_eaten_possible) |>
  mutate(size = intercept[sim] + slope[sim]*humans_eaten,
         resp = size + rnorm(n = n(), 0, sigma[sim])) -> prior_pred

ggplot(prior_pred, aes(x = humans_eaten, y = size)) +
  geom_line(aes(group = sim), linewidth = 0.1) +
  geom_hline(yintercept = 0, linetype = "dashed") +
  geom_point(aes(group = NULL), colour = "#3333B3", data = Alien)
```


LM — model fitting: prior predictive checking

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Before fitting the model, let's make sure our priors are not too vague


```
set.seed(123)
n_sims <- 1000

prior <- tibble(sim = seq_len(n_sims),
  intercept = rnorm(n_sims, mean = 50, sd = 5),
  slope = rnorm(n_sims, mean = 0, sd = 5),
  sigma = runif(n_sims, min = 0, max = 5))

humans_eaten_possible <- min(Alien$humans_eaten):max(Alien$humans_eaten)

expand_grid(prior, humans_eaten = humans_eaten_possible) |>
  mutate(size = intercept[sim] + slope[sim]*humans_eaten,
    resp = size + rnorm(n = n(), 0, sigma[sim])) -> prior_pred

ggplot(prior_pred, aes(x = humans_eaten, y = size)) +
  geom_line(aes(group = sim), linewidth = 0.1) +
  geom_hline(yintercept = 0, linetype = "dashed") +
  geom_point(aes(group = NULL), colour = "#3333B3", data = Alien)
```




- priors for intercept and slope look OK: wide enough to cover a wide range of relationships & not too wide
- scaling the data would improve things (see later)

LM — model fitting: prior predictive checking

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Before fitting the model, let's make sure our priors are not too vague

```
ggplot() +  
  geom_histogram(aes(x = resp, after_stat(density), fill = "from priors"),  
    data = prior_pred) +  
  geom_histogram(aes(x = size, after_stat(density), fill = "observed"),  
    alpha = 0.5, data = Alien) +  
  scale_fill_manual(values = c("#00CC99", "#3333B3"), name = "data")
```



- priors for sigma look OK too:
potentially the priors can produce data encompassing the real ones, without
risking producing silly alien sizes

LM — model fitting: prior predictive checking


- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Before fitting the model, let's make sure our priors are not too vague

```
ggplot() +  
  geom_histogram(aes(x = resp, after_stat(density), fill = "from priors"),  
                data = prior_pred) +  
  geom_histogram(aes(x = size, after_stat(density), fill = "observed"),  
                alpha = 0.5, data = Alien) +  
  scale_fill_manual(values = c("#00CC99", "#3333B3"), name = "data")
```

NB:

- one should not fine tune the priors based on the observed data since a dataset may be biased (risk of *overfitting*), but one should make sure that the priors are somewhat realistic
- how to best choose priors is an ongoing research question, and the most covered topic in Bayesian statistics (e.g. more objective techniques and databases are starting to emerge)
- if you have enough data, also long as priors are not too silly, they should have little effect on your posteriors (and you can check that)
- whatever you do, do describe what you do in your papers!




- priors for sigma look OK too: potentially the priors can produce data encompassing the real ones, without risking producing silly alien sizes

LM — model fitting: Bayesian approach with {nimble}

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

We can now fit the model with {nimble}

```
fit_alien_lm_nimb <- nimbleMCMC(code = code,  
                                constants = constants,  
                                data = list(size = Alien$size), # response goes here  
                                inits = inits, nchains = 4, setSeed = 1:4,  
                                nburnin = 4000) # amply sufficient here (brms used 1000)  
mcmc_trace(fit_alien_lm_nimb) # not mcmc_plot which only is for {brms}
```




LM — model fitting: Bayesian approach with {nimble}

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Here are the posterior distributions

```
mcmc_dens_overlay(fit_alien_lm_nimb)
```



LM — model fitting: Bayesian approach with {nimble}

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Here is the summary for the posteriors distributions

```
fit_alien_lm_nimb_d <- as_draws(fit_alien_lm_nimb)
summary(fit_alien_lm_nimb_d)
```

```
# A tibble: 3 x 10
  variable    mean median    sd   mad    q5   q95  rhat ess_bulk ess_tail
  <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>
1 intercept 49.3   49.3  1.20  1.17  47.4  51.3  1.00    3036.    6449.
2 sigma     1.98   1.89  0.502 0.439  1.33  2.93  1.00    3722.    4248.
3 slope     2.18   2.18  0.582 0.564  1.21  3.12  1.00    3042.    6209.
```

LM — model fitting: posterior predictive checking

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

After fitting the model, let's make sure our posterior are not off


```
posterior1 <- tibble(  
  intercept = fit_alien_lm_nimb_d[, 1, "intercept", drop = TRUE],  
  slope = fit_alien_lm_nimb_d[, 1, "slope", drop = TRUE],  
  sigma = fit_alien_lm_nimb_d[, 1, "sigma", drop = TRUE])  
  
expand_grid(posterior1, humans_eaten = Alien$humans_eaten) |>  
  mutate(error = rnorm(n = n(), mean = 0, sd = sigma),  
    pred = intercept + slope * humans_eaten + error) -> pred1  
  
ggplot() +  
  geom_density(aes(x = pred, after_stat(density),  
    colour = "from posteriors"), data = pred1) +  
  geom_density(aes(x = resp, after_stat(density),  
    colour = "from priors"), data = prior_pred,) +  
  geom_density(aes(x = size, after_stat(density),  
    colour = "observed"), data = Alien) +  
  scale_colour_manual(values = c("#FF5050", "#00CC99", "#3333B3"),  
    name = "data",  
    guide = guide_legend(position = "top")) +  
  xlim(limits = c(35, 75))
```

LM — model fitting: posterior predictive checking

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

After fitting the model, let's make sure our posterior are not off

```
posterior1 <- tibble(  
  intercept = fit_alien_lm_nimb_d[, 1, "intercept", drop = TRUE],  
  slope = fit_alien_lm_nimb_d[, 1, "slope", drop = TRUE],  
  sigma = fit_alien_lm_nimb_d[, 1, "sigma", drop = TRUE])  
  
expand_grid(posterior1, humans_eaten = Alien$humans_eaten) |>  
  mutate(error = rnorm(n = n(), mean = 0, sd = sigma),  
    pred = intercept + slope * humans_eaten + error) -> pred1  
  
ggplot() +  
  geom_density(aes(x = pred, after_stat(density),  
    colour = "from posteriors"), data = pred1) +  
  geom_density(aes(x = resp, after_stat(density),  
    colour = "from priors"), data = prior_pred,) +  
  geom_density(aes(x = size, after_stat(density),  
    colour = "observed"), data = Alien) +  
  scale_colour_manual(values = c("#FF5050", "#00CC99", "#3333B3"),  
    name = "data",  
    guide = guide_legend(position = "top")) +  
  xlim(limits = c(35, 75))
```




- this looks OK: the posterior predictions are much more concentrated than the prior ones, and close to the observations

LM — model fitting: posterior predictive checking

- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Let's check how correlated posteriors are

```
ggplot() +  
  geom_point(aes(y = intercept, x = slope, colour = "prior"),  
            alpha = 0.3, data = prior) +  
  geom_point(aes(y = intercept, x = slope, colour = "posterior"),  
            alpha = 0.3, data = posterior1) +  
  scale_colour_manual(  
    values = c("#00CC99", "#FF5050"),  
    name = "data",  
    guide = guide_legend(position = "top",  
                        override.aes = list(alpha = 1)))
```



- this looks OK: the posterior predictions are within the priors, which further validate the choice of our priors

LM — model fitting: posterior predictive checking


- in Bayesian statistics, *fitting the model* means using the data to characterise the posterior distributions

Let's check how correlated posteriors are

```
ggplot() +  
  geom_point(aes(y = intercept, x = slope, colour = "prior"),  
            alpha = 0.3, data = prior) +  
  geom_point(aes(y = intercept, x = slope, colour = "posterior"),  
            alpha = 0.3, data = posterior1) +  
  scale_colour_manual(  
    values = c("#00CC99", "#FF5050"),  
    name = "data",  
    guide = guide_legend(position = "top",  
                        override.aes = list(alpha = 1)))
```

NB: for real applications

- the predictor variables should be scaled to avoid strong correlations (see later)



- this looks OK: the posterior predictions are within the priors, which further validate the choice of our priors

LM — model fitting results: estimates compared

Let's compare the results

`lm()`

```
# A tibble: 3 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)    48.8        1.35    45.8     51.8
2 humans_eaten    2.37        0.623   0.987     3.76
3 sigma          1.76        1.39    1.01     6.36
```

`brms::brm()`

```
# A tibble: 3 x 5
  Parameter      Estimate Est.Error  Q2.5 Q97.5
  <chr>          <dbl>    <dbl>  <dbl> <dbl>
1 Intercept    48.8      1.51  45.8  51.8
2 humans_eaten  2.38      0.700  0.993  3.76
3 sigma        1.92      0.463  1.26  3.03
```

`nimble::nimbleMCMC()`

```
# A tibble: 3 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>      <dbl> <dbl>  <dbl>    <dbl>
1 intercept  49.3  1.20  47.0     51.7
2 slope      2.18  0.582  0.997     3.30
3 sigma      1.98  0.502  1.26     3.22
```

LM — model fitting results: estimates compared

Let's compare the results

```
lm()
```

```
# A tibble: 3 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)    48.8         1.35    45.8     51.8
2 humans_eaten    2.37         0.623   0.987     3.76
3 sigma          1.76         1.39    1.01     6.36
```

```
brms::brm()
```

```
# A tibble: 3 x 5
  Parameter      Estimate Est.Error  Q2.5 Q97.5
  <chr>          <dbl>      <dbl>  <dbl> <dbl>
1 Intercept    48.8        1.51  45.8  51.8
2 humans_eaten  2.38        0.700  0.993  3.76
3 sigma        1.92        0.463  1.26  3.03
```

```
nimble::nimbleMCMC()
```

```
# A tibble: 3 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>       <dbl> <dbl>  <dbl>    <dbl>
1 intercept  49.3   1.20  47.0     51.7
2 slope      2.18  0.582  0.997     3.30
3 sigma      1.98  0.502  1.26     3.22
```

- despite the small sample size (N=12) the estimates are in agreement and all intervals include ground truth values

LM — model fitting results: estimates compared

Let's compare the results

```
lm()
```

```
# A tibble: 3 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)    48.8        1.35    45.8     51.8
2 humans_eaten   2.37        0.623   0.987     3.76
3 sigma          1.76        1.39    1.01     6.36
```

```
brms::brm()
```

```
# A tibble: 3 x 5
  Parameter      Estimate Est.Error   Q2.5 Q97.5
  <chr>          <dbl>      <dbl>  <dbl> <dbl>
1 Intercept    48.8        1.51   45.8   51.8
2 humans_eaten  2.38        0.700   0.993   3.76
3 sigma        1.92        0.463   1.26   3.03
```

```
nimble::nimbleMCMC()
```

```
# A tibble: 3 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>      <dbl> <dbl>  <dbl>    <dbl>
1 intercept  49.3   1.20   47.0     51.7
2 slope      2.18  0.582  0.997     3.30
3 sigma      1.98  0.502  1.26     3.22
```

- despite the small sample size ($N=12$) the estimates are in agreement and all intervals include ground truth values
- we never obtained 50 (intercept), 1.5 (slope), or 2 (sigma) because a sample is unlikely to reflect perfectly the population-level values; only $N=\infty$ or talking to the goddess could get you those

LM — model fitting results: estimates compared

Let's compare the results

`lm()`

```
# A tibble: 3 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)    48.8        1.35    45.8     51.8
2 humans_eaten    2.37        0.623   0.987     3.76
3 sigma          1.76        1.39    1.01     6.36
```

`brms::brm()`

```
# A tibble: 3 x 5
  Parameter      Estimate Est.Error  Q2.5 Q97.5
  <chr>          <dbl>      <dbl>  <dbl> <dbl>
1 Intercept    48.8        1.51   45.8   51.8
2 humans_eaten  2.38        0.700   0.993   3.76
3 sigma        1.92        0.463   1.26   3.03
```

`nimble::nimbleMCMC()`

```
# A tibble: 3 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>      <dbl> <dbl>  <dbl>    <dbl>
1 intercept  49.3   1.20   47.0     51.7
2 slope      2.18  0.582  0.997     3.30
3 sigma      1.98  0.502  1.26     3.22
```

- despite the small sample size (N=12) the estimates are in agreement and all intervals include ground truth values
- we never obtained 50 (intercept), 1.5 (slope), or 2 (sigma) because a sample is unlikely to reflect perfectly the population-level values; only $N=\infty$ or talking to the goddess could get you those

Geeky notes:

- in `lm()` $\hat{\sigma}$ is usually not a parameter that is estimated directly, it is deduced once the intercept and slope(s) have been fitted (behind the scenes I struggled to compute the SE and CI for it)

LM — model fitting results: estimates compared

Let's compare the results

`lm()`

```
# A tibble: 3 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)    48.8        1.35    45.8     51.8
2 humans_eaten   2.37        0.623   0.987     3.76
3 sigma          1.76        1.39    1.01     6.36
```

`brms::brm()`

```
# A tibble: 3 x 5
  Parameter      Estimate Est.Error  Q2.5 Q97.5
  <chr>          <dbl>    <dbl>  <dbl> <dbl>
1 Intercept    48.8      1.51  45.8  51.8
2 humans_eaten  2.38      0.700  0.993  3.76
3 sigma        1.92      0.463  1.26  3.03
```

`nimble::nimbleMCMC()`

```
# A tibble: 3 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>      <dbl> <dbl>  <dbl>    <dbl>
1 intercept  49.3  1.20  47.0     51.7
2 slope      2.18 0.582  0.997     3.30
3 sigma      1.98 0.502  1.26     3.22
```

- despite the small sample size ($N=12$) the estimates are in agreement and all intervals include ground truth values
- we never obtained 50 (intercept), 1.5 (slope), or 2 (sigma) because a sample is unlikely to reflect perfectly the population-level values; only $N=\infty$ or talking to the goddess could get you those

Geeky notes:

- in `lm()` $\hat{\sigma}$ is usually not a parameter that is estimated directly, it is deduced once the intercept and slope(s) have been fitted (behind the scenes I struggled to compute the SE and CI for it)
- by default nimble returns q5 and q95 (not q2.5 and q97.5), so I had to request this manually (see `?summarise_draws`)

Practice 2

1. explore the dataset `penguins` (readily available in R)
2. do check its documentation `?penguins`
3. using this dataset, fit a linear model with:
 - `body_mass` as the response variables
 - `flipper_len` as predictors

Interpreting coefficients

Interpreting coefficients — introduction

Interpreting the meaning of *regression coefficients* (= *parameter estimates*) correctly in statistical models is the most critical part given that:

- coefficients mediate the relationship between the statistics and the biology

Interpreting coefficients — introduction

Interpreting the meaning of *regression coefficients* (= *parameter estimates*) correctly in statistical models is the most critical part given that:

- coefficients mediate the relationship between the statistics and the biology
- tests, SE, confidence intervals and prediction intervals, all focus on these coefficients

Interpreting coefficients — introduction

Interpreting the meaning of *regression coefficients* (= *parameter estimates*) correctly in statistical models is the most critical part given that:

- coefficients mediate the relationship between the statistics and the biology
- tests, SE, confidence intervals and prediction intervals, all focus on these coefficients
- different software and packages may influence aspects (contrasts, scale) that impact the meaning of the coefficients

Interpreting coefficients — introduction

Interpreting the meaning of *regression coefficients* (= *parameter estimates*) correctly in statistical models is the most critical part given that:

- coefficients mediate the relationship between the statistics and the biology
- tests, SE, confidence intervals and prediction intervals, all focus on these coefficients
- different software and packages may influence aspects (contrasts, scale) that impact the meaning of the coefficients
- most scientists fail to do so correctly, jeopardising the conclusions drawn in many publications

Interpreting coefficients — introduction

Interpreting the meaning of *regression coefficients* (= *parameter estimates*) correctly in statistical models is the most critical part given that:

- coefficients mediate the relationship between the statistics and the biology
- tests, SE, confidence intervals and prediction intervals, all focus on these coefficients
- different software and packages may influence aspects (contrasts, scale) that impact the meaning of the coefficients
- most scientists fail to do so correctly, jeopardising the conclusions drawn in many publications

NB:

- it requires some focus, but no more than middle school mathematics knowledge

Interpreting coefficients — data preparation: clean up

Let's create a single clean dataset with all the variables we will need for all models

```
penguins |>
  as_tibble() |>
  select(-starts_with("bill")) |>
  drop_na() |>
  mutate(flipper_len_z = scale(flipper_len)[,1], .after = flipper_len) |>
  mutate(body_mass_z = scale(body_mass)[,1], .after = body_mass) -> penguins_clean

penguins_clean
```

A tibble: 333 x 8


	species	island	flipper_len	flipper_len_z	body_mass	body_mass_z	sex	year
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>	<fct>	<int>
1	Adelie	Torgersen	181	-1.42	3750	-0.568	male	2007
2	Adelie	Torgersen	186	-1.07	3800	-0.506	female	2007
3	Adelie	Torgersen	195	-0.426	3250	-1.19	female	2007
4	Adelie	Torgersen	193	-0.568	3450	-0.940	female	2007
5	Adelie	Torgersen	190	-0.782	3650	-0.692	male	2007

i 328 more rows

NB: `scale()` creates z-scores (it removes the mean and divide by the SD)

Interpreting coefficients — data preparation: plot

```
ggplot(penguins_clean) +  
  aes(y = body_mass, x = flipper_len, colour = species, shape = sex) +  
  geom_point(alpha = 0.5) +  
  labs(y = "Body mass (gr)", x = "Flipper length (mm)")
```



Interpreting coefficients — intercept and slopes

How to interpret the regression coefficients?

```
fit_flipper <- lm(body_mass ~ flipper_len, data = penguins_clean)
coef(fit_flipper)
```

```
(Intercept) flipper_len
-5872.09268    50.15327
```

Interpreting coefficients — intercept and slopes

How to interpret the regression coefficients?

```
fit_flipper <- lm(body_mass ~ flipper_len, data = penguins_clean)
coef(fit_flipper)
```

```
(Intercept) flipper_len
-5872.09268    50.15327
```

- the *intercept* is the prediction when all regressors are at 0, so here it corresponds to the average body mass of a penguin that would have no flippers (i.e. a nonsensical *extrapolation* in this case)

```
predict(fit_flipper, newdata = tibble(flipper_len = 0))
```

```
1
-5872.093
```

Interpreting coefficients — intercept and slopes

How to interpret the regression coefficients?

```
fit_flipper <- lm(body_mass ~ flipper_len, data = penguins_clean)
coef(fit_flipper)
```

```
(Intercept) flipper_len
-5872.09268    50.15327
```

- the *slope* gives the change in predicted values when its associated regressor increases by 1 unit, so here it means that the model predicts that penguins with flippers 1mm longer than others are on average 50gr heavier than such others

```
pred_diff <- predict(fit_flipper, newdata = tibble(flipper_len = c(200, 201)))
pred_diff
```

```
      1      2
4158.561 4208.714
```

```
pred_diff[2] - pred_diff[1] # or simply diff(pred_diff)
```

```
      2
50.15327
```

Interpreting coefficients — scaling predictors

How to interpret the regression coefficients when predictors are scaled?

```
fit_flipper_z <- lm(body_mass ~ flipper_len_z, data = penguins_clean)
coef(fit_flipper_z)
```

```
(Intercept) flipper_len_z
  4207.0571      702.9364
```

Interpreting coefficients — scaling predictors

How to interpret the regression coefficients when predictors are scaled?

```
fit_flipper_z <- lm(body_mass ~ flipper_len_z, data = penguins_clean)
coef(fit_flipper_z)
```

```
(Intercept) flipper_len_z
  4207.0571      702.9364
```

NB:

- *scaling* means removing the mean (i.e. *centering*) and dividing by the SD (i.e. *standardising*)
- scaling both the response variable and the predictors is a recommended practice when using {nimble}, {stan}, etc... since it makes the definition of priors more consistent across studies; it also improves the efficiency of some MCMC samplers since it decreases correlations between posteriors
- some packages do similar things internally (e.g. {brms}), while doing themselves the necessary *back transformation* for allowing interpretations at the original scales

Interpreting coefficients — scaling predictors

How to interpret the regression coefficients when predictors are scaled?

```
fit_flipper_z <- lm(body_mass ~ flipper_len_z, data = penguins_clean)
coef(fit_flipper_z)
```

```
(Intercept) flipper_len_z
  4207.0571      702.9364
```

- the *intercept* remains the prediction when all regressors are at 0, but `flipper_len = 0` meant a flipper length of 0mm, while `flipper_len_z = 0` means an average flipper length; so here the intercept corresponds to the average body mass of a penguin that would have flippers of 201mm (i.e. no longer nonsensical)

Interpreting coefficients — scaling predictors

How to interpret the regression coefficients when predictors are scaled?

```
fit_flipper_z <- lm(body_mass ~ flipper_len_z, data = penguins_clean)
coef(fit_flipper_z)
```

```
(Intercept) flipper_len_z
  4207.0571      702.9364
```

- the *slope* remains the change in predicted values when its associated regressor increases by 1 unit, but here 1 unit is not longer 1mm but 1 SD in flipper length; so here it means that the model predicts that penguins with flippers 1 SD longer than others (14mm) are on average 703gr heavier than such others

back transformation

It is possible to convert the coefficients to get back to the original (untransformed) scale!

```
coef(fit_flipper_z)["flipper_len_z"] / sd(penguins_clean$flipper_len)
```

```
flipper_len_z
  50.15327
```

Interpreting coefficients — scaling the response

How to interpret the regression coefficients when the response is scaled?

```
fit_flipper_bmz <- lm(body_mass_z ~ flipper_len, data = penguins_clean)
coef(fit_flipper_bmz)
```

```
(Intercept) flipper_len
-12.5173273    0.0622855
```

- again, the *intercept* remains the prediction when all regressors are at 0, and the *slope* remains the change in predicted values when its associated regressor increases by 1 unit, but this time the predicted values are expressed in numbers of SDs above the mean (i.e. z-scores) rather than grams

back transformation

```
coef(fit_flipper_bmz)["(Intercept)"] * sd(penguins_clean$body_mass) + mean(penguins_clean$body_mass)
```

```
(Intercept)
-5872.093
```

```
coef(fit_flipper_bmz)["flipper_len"] * sd(penguins_clean$body_mass)
```

```
flipper_len
50.15327
```


Practice 3

How to interpret the regression coefficients when the response and predictors are scaled?

```
fit_flipper_zz <- lm(body_mass_z ~ flipper_len_z, data = penguins_clean)
coef(fit_flipper_zz)
```

```
(Intercept) flipper_len_z  
-2.689131e-16  8.729789e-01
```

1. interpret the slope on its new scale
2. interpret the intercept on its new scale
3. apply the required back transformation to perform the interpretation of the slope on the original scale
4. think about how you would define the priors in this case

Solution

How to interpret the regression coefficients when the response and predictors are scaled?

```
fit_flipper_zz <- lm(body_mass_z ~ flipper_len_z, data = penguins_clean)
coef(fit_flipper_zz)
```

```
(Intercept) flipper_len_z
-2.689131e-16  8.729789e-01
```

1. the slope describes the increase in SD of body mass for an increase of 1 SD in flipper length
2. the intercept is the mean of the body mass z-scores for a penguin with average flipper length, which is why it is 0 (a group of penguins with an average flipper length is expected to have an average body mass)
3. back transformation of the slope:

```
(coef(fit_flipper_zz)["flipper_len_z"] / sd(penguins_clean$flipper_len)) * sd(penguins_clean$body_mass)
```

```
flipper_len_z
50.15327
```

Solution

4. something like $\mathcal{N}(0, \sigma = 2)$ for both the intercept and slope priors seems appropriate (& close to {brms} default settings)


```
set.seed(123)
n_sims <- 100

prior <- tibble(sim = seq_len(n_sims),
                intercept = rnorm(n_sims, mean = 0, sd = 2),
                slope = rnorm(n_sims, mean = 0, sd = 2))

flipper_len_z_possible <- -3:3

expand_grid(prior, flipper_len_z = flipper_len_z_possible) |>
  mutate(
    body_mass_z = intercept[sim] + slope[sim]*flipper_len_z) -> prior_pred

ggplot(prior_pred, aes(flipper_len_z, body_mass_z)) +
  geom_line(aes(group = sim), linewidth = 0.1) +
  geom_hline(yintercept = 0, linetype = "dashed") +
  geom_point(aes(group = NULL), data = penguins_clean)
```



Interpreting coefficients — qualitative predictors

How to interpret the regression coefficients?

```
fit_sex <- lm(body_mass ~ sex, data = penguins_clean)
coef(fit_sex)
```

(Intercept)	sexmale
3862.2727	683.4118

Interpreting coefficients — qualitative predictors

How to interpret the regression coefficients?

```
fit_sex <- lm(body_mass ~ sex, data = penguins_clean)
coef(fit_sex)
```

(Intercept)	sexmale
3862.2727	683.4118

Any idea?

Interpreting coefficients — qualitative predictors

How to interpret the regression coefficients?

```
fit_sex <- lm(body_mass ~ sex, data = penguins_clean)
coef(fit_sex)
```

(Intercept)	sexmale
3862.2727	683.4118

Any idea?

NB: it is usually best to avoid lucky guesses in this context

Interpreting coefficients — model matrix interlude

It is more useful to check how the data are converted into the *model matrix*:

```
penguins_clean # data
```

```
# A tibble: 333 x 8
```

	species	island	flipper_len	flipper_len_z	body_mass	body_mass_z	sex	year
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>	<fct>	<int>
1	Adelie	Torgersen	181	-1.42	3750	-0.568	male	2007
2	Adelie	Torgersen	186	-1.07	3800	-0.506	female	2007
3	Adelie	Torgersen	195	-0.426	3250	-1.19	female	2007
4	Adelie	Torgersen	193	-0.568	3450	-0.940	female	2007
5	Adelie	Torgersen	190	-0.782	3650	-0.692	male	2007

```
# i 328 more rows
```

```
head(model.matrix(fit_sex))
```

	(Intercept)	sexmale
1	1	1
2	1	0
3	1	0
4	1	0
5	1	1
6	1	0

Interpreting coefficients — model matrix interlude

It is more useful to check how the data are converted into the *model matrix*:

```
penguins_clean # data
```

```
# A tibble: 333 x 8
  species island flipper_len flipper_len_z body_mass body_mass_z sex    year
  <fct>   <fct>      <int>      <dbl>      <int>      <dbl> <fct> <int>
1 Adelie Torgersen    181      -1.42      3750      -0.568 male   2007
2 Adelie Torgersen    186      -1.07      3800      -0.506 female 2007
3 Adelie Torgersen    195      -0.426     3250      -1.19  female 2007
4 Adelie Torgersen    193      -0.568     3450      -0.940 female 2007
5 Adelie Torgersen    190      -0.782     3650      -0.692 male   2007
# i 328 more rows
```

```
head(model.matrix(fit_sex))
```

```
(Intercept) sexmale
1           1       1
2           1       0
3           1       0
4           1       0
5           1       1
6           1       0
```

- the model matrix is multiplied by the regression coefficients, so the coefficient `sexmale` is multiplied by 1 when the row corresponds to a male, and by 0 when it corresponds to a female

Interpreting coefficients — model matrix interlude

It is more useful to check how the data are converted into the *model matrix*:

```
penguins_clean # data
```

```
# A tibble: 333 x 8
  species island flipper_len flipper_len_z body_mass body_mass_z sex    year
  <fct>   <fct>         <int>         <dbl>         <int>         <dbl> <fct> <int>
1 Adelie  Torgersen         181         -1.42          3750        -0.568 male   2007
2 Adelie  Torgersen         186         -1.07          3800        -0.506 female 2007
3 Adelie  Torgersen         195         -0.426         3250        -1.19  female 2007
4 Adelie  Torgersen         193         -0.568         3450        -0.940 female 2007
5 Adelie  Torgersen         190         -0.782         3650        -0.692 male   2007
# i 328 more rows
```

```
head(model.matrix(fit_sex))
```

	(Intercept)	sexmale
1	1	1
2	1	0
3	1	0
4	1	0
5	1	1
6	1	0

- the model matrix is multiplied by the regression coefficients, so the coefficient `sexmale` is multiplied by 1 when the row corresponds to a male, and by 0 when it corresponds to a female
- by default, R chooses a reference level (by default, the first in alphabetical order) and expresses other effects in relation to that reference, a method called *treatment contrast*

Interpreting coefficients — qualitative predictors

How to interpret the regression coefficients?

```
coef(fit_sex)
```

(Intercept)	sexmale
3862.2727	683.4118

predicted value_{*i*} = (Intercept) × 1 + sexmale × (0 if *i* is female or 1 if *i* is male)

- the *intercept* remains the prediction when all regressors are at 0, so here the average body mass of a penguin who is *female*
- the value called *sexmale* tells you by how much the predicted values increase when the column of the design matrix called *sexmale* changes from 0 to 1 → it is thus the average increase in body mass of male penguins *compared to* female penguins

When in doubt, do compare predictions!

```
pred_sex <- predict(fit_sex, newdata = tibble(sex = c("female", "male")))
pred_sex
```

1	2
3862.273	4545.685

Practice 4

How to interpret the regression coefficients with an interaction between a qualitative and quantitative predictor?

```
fit_flipper_int <- lm(body_mass ~ flipper_len*sex, data = penguins_clean)
coef(fit_flipper_int)
```

(Intercept)	flipper_len	sexmale	flipper_len:sexmale
-5443.9607499	47.1527260	406.8015395	-0.2942184

1. check the design matrix to understand the regressor used by `flipper_len:sexmale`
2. predict by hand:
 - the average body mass of a female penguin with a flipper length of 200 mm
 - the average body mass of a male penguin with a flipper length of 200 mm
3. interpret the meaning of all 4 coefficients
4. does body mass increase with flipper length in males?

Solution

```
coef(fit_flipper_int)
```

(Intercept)	flipper_len	sexmale	flipper_len:sexmale
-5443.9607499	47.1527260	406.8015395	-0.2942184

1. the design matrix tells us that `flipper_len:sexmale` is 0 for a female and the flipper length when the penguin is male

```
head(model.matrix(fit_flipper_int), n = 2)
```

	(Intercept)	flipper_len	sexmale	flipper_len:sexmale
1	1	181	1	181
2	1	186	0	0

```
head(penguins_clean, n = 2)
```

A tibble: 2 x 8

	species	island	flipper_len	flipper_len_z	body_mass	body_mass_z	sex	year
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>	<fct>	<int>
1	Adelie	Torgersen	181	-1.42	3750	-0.568	male	2007
2	Adelie	Torgersen	186	-1.07	3800	-0.506	female	2007

Solution

```
coef(fit_flipper_int)
```

(Intercept)	flipper_len	sexmale	flipper_len:sexmale
-5443.9607499	47.1527260	406.8015395	-0.2942184

2. predictions can be computed as follows:

predicted value_{*i*} = (Intercept) × 1 + flipper_len × flipper_len_{*i*} + sexmale ×
(0 if *i* is female or 1 if *i* is male) + flipper_len:sexmale × (0 if *i* is female or flipper_len_{*i*} if *i* is male)

- the average body mass of a female penguin with a flipper length of 200 mm

$$= -5443.96 + 47.15 \times 200 + 406.80 \times 0 - 0.29 \times 0 \sim 3986$$

Solution

```
coef(fit_flipper_int)
```

(Intercept)	flipper_len	sexmale	flipper_len:sexmale
-5443.9607499	47.1527260	406.8015395	-0.2942184

2. predictions can be computed as follows:

predicted value_{*i*} = (Intercept) × 1 + flipper_len × flipper_len_{*i*} + sexmale ×
(0 if *i* is female or 1 if *i* is male) + flipper_len:sexmale × (0 if *i* is female or flipper_len_{*i*} if *i* is male)

- the average body mass of a female penguin with a flipper length of 200 mm

$$= -5443.96 + 47.15 \times 200 + 406.80 \times 0 - 0.29 \times 0 \sim 3986$$

- the average body mass of a male penguin with a flipper length of 200 mm

$$= -5443.96 + 47.15 \times 200 + 406.80 \times 1 - 0.29 \times 200 = 3986 \sim 4335$$

Exact solution (without rounding)

```
predict(fit_flipper_int, newdata = tibble(sex = c("female", "male"), flipper_len = 200))
```

1 2

3986.584 4334.542

Solution

```
coef(fit_flipper_int)
```

(Intercept)	flipper_len	sexmale	flipper_len:sexmale
-5443.9607499	47.1527260	406.8015395	-0.2942184

3. interpretation

- `(Intercept)` = average mean body mass of females with flippers of 0mm
- `flipper_len` = average increase in body mass for females when flipper length is increased by 1mm
- `sexmale` = average increase in male body mass **compared to female body mass** for penguins with flippers of 0mm
- `flipper_len:sexmale` = average increase in male body mass **compared to the increase in females** when flipper length is increased by 1mm

Solution

```
coef(fit_flipper_int)
```

(Intercept)	flipper_len	sexmale	flipper_len:sexmale
-5443.9607499	47.1527260	406.8015395	-0.2942184

4. body mass does increase with flipper length in males since the effect of flipper length in males is $47.15 - 0.29 (>> 0)$


```
predict(fit_flipper_int, newdata = tibble(sex = "male", flipper_len = c(200, 201)))
```

1	2
4334.542	4381.401

Interpreting coefficients — plotting predictions

Computing and plotting prediction can be difficult, but the package `{modelbased}` from the `{easystats}` universe (easystats.github.io/easystats) makes this easy for many types of models

```
library(modelbased)
estimate_relation(fit_flipper_int) |>
  plot(show_data = TRUE)
```



Interpreting coefficients — conditional and marginal predictions

How to compute/illustrate model predictions for a subset of predictors?

```
fit_flipper_3var <- lm(body_mass ~ flipper_len + sex + species, data = penguins_clean)
```

Example

you want to compute/illustrate the relationship between `body_mass` & `flipper_len` from the fitted model above, what to do with `sex` & `species`?

¹structural correlations are caused by an unbalanced sampling design

Interpreting coefficients — conditional and marginal predictions

How to compute/illustrate model predictions for a subset of predictors?

```
fit_flipper_3var <- lm(body_mass ~ flipper_len + sex + species, data = penguins_clean)
```

Example

you want to compute/illustrate the relationship between `body_mass` & `flipper_len` from the fitted model above, what to do with `sex` & `species`?

- **conditional predictions**: the software fixes each non-focal predictor at a specific value (e.g. the mean for quantitative variables & the most frequent category for the qualitative variables) → easiest, but choice of reference can be arbitrary and the computation is even slightly wrong in models involving non-linear transformation(s) (due to *Jensen's inequality*)

¹structural correlations are caused by an unbalanced sampling design

Interpreting coefficients — conditional and marginal predictions

How to compute/illustrate model predictions for a subset of predictors?

```
fit_flipper_3var <- lm(body_mass ~ flipper_len + sex + species, data = penguins_clean)
```

Example

you want to compute/illustrate the relationship between `body_mass` & `flipper_len` from the fitted model above, what to do with `sex` & `species`?

- **conditional predictions:** the software fixes each non-focal predictor at a specific value (e.g. the mean for quantitative variables & the most frequent category for the qualitative variables) → easiest, but choice of reference can be arbitrary and the computation is even slightly wrong in models involving non-linear transformation(s) (due to *Jensen's inequality*)
- **marginal predictions:** the software computes predictions for different values/levels of the focal and non-focal predictors and then averages those predictions across the non-focal predictors → often better, but it is more difficult to grasp since different alternatives exist to compute the *reference grid* (depending on whether one wants to capture or erase the actual or structural correlations between predictors¹)

¹structural correlations are caused by an unbalanced sampling design

Interpreting coefficients — conditional and marginal predictions

Conditional predictions of the effect of `flipper_len` on `body_mass`

```
cond_pred <- estimate_relation(fit_flipper_3var,  
                              by = "flipper_len")
```

`cond_pred`

Model-based Predictions


flipper_len	Predicted	SE	95% CI
172.00	3078.47	51.04	[2978.05, 3178.88]
178.56	3209.75	37.35	[3136.28, 3283.22]
185.11	3341.01	29.68	[3282.64, 3399.39]
191.67	3472.30	32.59	[3408.18, 3536.42]
198.22	3603.56	44.05	[3516.91, 3690.21]
204.78	3734.84	59.28	[3618.23, 3851.46]
211.33	3866.11	76.05	[3716.50, 4015.72]
217.89	3997.39	93.54	[3813.38, 4181.41]
224.44	4128.65	111.41	[3909.49, 4347.82]
231.00	4259.94	129.50	[4005.19, 4514.69]

Variable predicted: `body_mass`

Predictors modulated: `flipper_len`

Predictors controlled: `sex (female)`, `species (Adelie)`

```
plot(cond_pred, show_data = TRUE)
```



Interpreting coefficients — conditional and marginal predictions

Conditional predictions of the effect of `flipper_len` on `body_mass`

```
cond_pred <- estimate_relation(fit_flipper_3var,  
                              by = "flipper_len")
```

`cond_pred`

Model-based Predictions


flipper_len	Predicted	SE	95% CI
172.00	3078.47	51.04	[2978.05, 3178.88]
178.56	3209.75	37.35	[3136.28, 3283.22]
185.11	3341.01	29.68	[3282.64, 3399.39]
191.67	3472.30	32.59	[3408.18, 3536.42]
198.22	3603.56	44.05	[3516.91, 3690.21]
204.78	3734.84	59.28	[3618.23, 3851.46]
211.33	3866.11	76.05	[3716.50, 4015.72]
217.89	3997.39	93.54	[3813.38, 4181.41]
224.44	4128.65	111.41	[3909.49, 4347.82]
231.00	4259.94	129.50	[4005.19, 4514.69]

Variable predicted: `body_mass`

Predictors modulated: `flipper_len`

Predictors controlled: `sex (female)`, `species (Adelie)`

```
plot(cond_pred, show_data = TRUE)
```



NB: here, the conditional predictions are given for female Adelie penguins; in this particular model (but not generally), another reference (e.g. male Gentoo penguins) would only shift the regression line up or down, but it would not modify the slope since no interactions are considered and the effects are linear


Interpreting coefficients — conditional and marginal predictions

Marginal predictions of the effect of `flipper_len` on `body_mass` (“average” method)

```
marg_avg_pred <- estimate_means(  
  fit_flipper_3var,  
  by = list(flipper_len = penguins_clean$flipper_len),  
  estimate = "average")  
  
as_tibble(marg_avg_pred)
```

```
# A tibble: 54 x 7  
  flipper_len Mean      SE CI_low CI_high      t      df  
    <int> <dbl> <dbl> <dbl> <dbl> <dbl> <int>  
1      172 3078.  51.0  2978.  3179.  60.3  328  
2      174 3119.  46.5  3027.  3210.  67.1  328  
3      176 3159.  42.2  3075.  3242.  74.8  328  
4      178 3309.  40.2  3230.  3388.  82.4  328  
5      180 3636.  42.4  3553.  3720.  85.8  328  
# i 49 more rows
```

```
plot(marg_avg_pred, show_data = TRUE)
```



NB: here, the marginal prediction is the average of the predicted values for all sampled penguins with the same flipper length and which may differ in their sex or species
→ preserves correlations between predictors & preserves the distributions of non-focal predictors (i.e. sex & species sampled more weigh more)

Interpreting coefficients — conditional and marginal predictions

Marginal predictions of the effect of `flipper_len` on `body_mass` (“average” method)

```
marg_avg_pred <- estimate_means(  
  fit_flipper_3var,  
  by = list(flipper_len = penguins_clean$flipper_len),  
  estimate = "average")  
  
as_tibble(marg_avg_pred)
```

```
# A tibble: 54 x 7  
  flipper_len Mean      SE CI_low CI_high      t      df  
    <int> <dbl> <dbl> <dbl> <dbl> <dbl> <int>  
1      172 3078.  51.0  2978.  3179.  60.3  328  
2      174 3119.  46.5  3027.  3210.  67.1  328  
3      176 3159.  42.2  3075.  3242.  74.8  328  
4      178 3309.  40.2  3230.  3388.  82.4  328  
5      180 3636.  42.4  3553.  3720.  85.8  328  
# i 49 more rows
```

```
penguins_clean |>  
  select(flipper_len, sex, species) -> grid_avg
```

```
grid_avg |>  
  mutate(pred = predict(  
    fit_flipper_3var,  
    newdata = grid_avg)) -> grid_avg
```

```
grid_avg |>  
  summarise(Mean = mean(pred), .by = "flipper_len") |>  
  arrange(flipper_len)
```

```
# A tibble: 54 x 2  
  flipper_len Mean  
    <int> <dbl>  
1      172 3078.  
2      174 3119.  
3      176 3159.  
4      178 3309.  
5      180 3636.  
# i 49 more rows
```



Interpreting coefficients — conditional and marginal predictions

Marginal predictions of the effect of `flipper_len` on `body_mass` (“typical” method)

```
marg_typic_pred <- estimate_means(  
  fit_flipper_3var,  
  by = "flipper_len",  
  estimate = "typical")  
  
as_tibble(marg_typic_pred)
```

```
# A tibble: 10 x 7  
  flipper_len Mean      SE CI_low CI_high      t    df  
    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>  
1      172 3593.  84.3 3427.  3759.  42.6  328  
2      179. 3724.  66.2 3594.  3855.  56.3  328  
3      185. 3856.  48.4 3761.  3951.  79.7  328  
4      192. 3987.  31.6 3925.  4049. 126.  328  
5      198. 4118.  18.8 4081.  4155. 219.  328  
# i 5 more rows
```

```
plot(marg_typic_pred, show_data = TRUE)
```



NB: here, the marginal prediction is the average of the predicted values considering that any given flipper length could be found in either sex and in any species with the same probability → breaks down correlations between predictors & flattens the distributions of non-focal predictors (i.e. both sex and all species exert the same weight, as in a perfectly balanced experimental design)

Interpreting coefficients — conditional and marginal predictions

Marginal predictions of the effect of `flipper_len` on `body_mass` (“typical” method)

```
marg_typic_pred <- estimate_means(  
  fit_flipper_3var,  
  by = "flipper_len",  
  estimate = "typical")  
  
as_tibble(marg_typic_pred)
```

```
# A tibble: 10 x 7  
  flipper_len Mean      SE CI_low CI_high      t      df  
    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>  
1      172 3593.  84.3 3427. 3759.  42.6  328  
2      179. 3724.  66.2 3594. 3855.  56.3  328  
3      185. 3856.  48.4 3761. 3951.  79.7  328  
4      192. 3987.  31.6 3925. 4049. 126.  328  
5      198. 4118.  18.8 4081. 4155. 219.  328  
# i 5 more rows
```

```
expand_grid(  
  flipper_len = seq(172, 231, length = 10),  
  sex = unique(penguins_clean$sex),  
  species = unique(penguins_clean$species)) -> grid_typic  
  
grid_typic |>  
  mutate(pred = predict(  
    fit_flipper_3var,  
    newdata = grid_typic)) -> grid_typic  
  
grid_typic |>  
  summarise(Mean = mean(pred), .by = "flipper_len") |>  
  arrange(flipper_len)
```

```
# A tibble: 10 x 2  
  flipper_len Mean  
    <dbl> <dbl>  
1      172 3593.  
2      179. 3724.  
3      185. 3856.  
4      192. 3987.  
5      198. 4118.  
# i 5 more rows
```


Interpreting coefficients — conditional and marginal predictions

Marginal predictions of the effect of `flipper_len` on `body_mass` (“population” method)

```
marg_pop_pred <- estimate_means(  
  fit_flipper_3var,  
  by = "flipper_len",  
  estimate = "population")  
  
as_tibble(marg_pop_pred)
```

```
# A tibble: 10 x 7  
  flipper_len Mean      SE CI_low CI_high      t    df  
    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>  
1      172 3627.  84.0 3462.  3792.  43.2  328  
2      179 3758.  65.8 3629.  3888.  57.1  328  
3      185 3890.  47.9 3795.  3984.  81.1  328  
4      192 4021.  31.0 3960.  4082. 130.  328  
5      198 4152.  18.0 4117.  4187. 231.  328  
# i 5 more rows
```

```
plot(marg_pop_pred, show_data = TRUE)
```



NB: here, the marginal prediction is the average of the predicted values considering that any given flipper length could be found in either sex and in any species with probability reflecting the sampled data → breaks down correlations between predictors, but preserve the distributions of non-focal predictors (i.e. sex & species sampled more weigh more)

Interpreting coefficients — conditional and marginal predictions

Marginal predictions of the effect of `flipper_len` on `body_mass` (“population” method)

```
marg_pop_pred <- estimate_means(  
  fit_flipper_3var,  
  by = "flipper_len",  
  estimate = "population")
```

```
as_tibble(marg_pop_pred)
```

```
# A tibble: 10 x 7  
  flipper_len Mean    SE CI_low CI_high    t    df  
    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>  
1      172 3627.  84.0 3462.  3792.  43.2  328  
2      179. 3758.  65.8 3629.  3888.  57.1  328  
3      185. 3890.  47.9 3795.  3984.  81.1  328  
4      192. 4021.  31.0 3960.  4082. 130.  328  
5      198. 4152.  18.0 4117.  4187. 231.  328  
# i 5 more rows
```

```
expand_grid(  
  flipper_len = seq(172, 231, length = 10),  
  sex = penguins_clean$sex,  
  species = penguins_clean$species) -> grid_pop
```

```
grid_pop |>  
  mutate(pred = predict(  
    fit_flipper_3var,  
    newdata = grid_pop)) -> grid_pop
```

```
grid_pop |>  
  summarise(Mean = mean(pred), .by = "flipper_len") |>  
  arrange(flipper_len)
```

```
# A tibble: 10 x 2  
  flipper_len Mean  
    <dbl> <dbl>  
1      172 3627.  
2      179. 3758.  
3      185. 3890.  
4      192. 4021.  
5      198. 4152.  
# i 5 more rows
```

Generalised Linear Models (GLM)

GLM — why the need?

Classic linear models (LM) are useful to fit many phenomena, including many for which the relationship between a predictor and the response is not linear (e.g. polynomials), but they share the limitation that they are designed to approximate a data generating process that is intrinsically Gaussian (remember that the error is assumed normally distributed)

LM won't work well for:

- binary data (e.g. the presence/absence of a species)
- proportion (e.g. the hatching success in a nest), especially if values get close to 0 or 1
- count data (e.g. the number of species caught on camera), especially if counts are low
- time to event data (e.g. time to dispersion), especially if times are short

For these situations GLM are the first choice!

NB: logistic regression & poisson regressions are examples of GLM

GLM — notations of a GLM: simple notation

$$\eta_i = \beta_0 + \beta_1 \times x_{1,i} + \beta_2 \times x_{2,i} + \dots + \beta_p \times x_{p,i}$$

with

- $E(Y) = \mu = g^{-1}(\eta)$
- $\text{Var}(Y) = \phi V(\mu)$

We call:

- η the linear predictor
- g the link function (g^{-1} is the inverse link function)
- V the variance function
- ϕ is the dispersion parameter

NB:

- it is very much like an LM but there are new elements
- in GLM, if $\phi = 1$ (the usual case) the variance in the response is directly a function of the mean and thus of η , so no need for σ in this case!

GLM — data generating process: data simulation

Let's reenter **God(dess) mode** → we pretend we know how her world generates data

GLM — data generating process: data simulation

Let's reenter **God(dess) mode** → we pretend we know how her world generates data

Emerald Lion informed us that in her world aliens depart to another planet following this rule:

$$P(\text{departure}_i) = \frac{1}{1 + \exp^{-(2.197 + 1.5 \times \text{humans_eaten}_i)}}$$



NB: the function $\frac{1}{1 + e^{-x}}$ is called the standard logistic function and is readily available in R via `plogis(x)`

GLM — data generating process: data simulation

Let's reenter **God(dess) mode** → we pretend we know how her world generates data

Emerald Lion informed us that in her world aliens depart to another planet following this rule:

$$P(\text{departure}_i) = \frac{1}{1 + \exp^{-(-2.197 + 1.5 \times \text{humans_eaten}_i)}}$$

We can thus generate data on 35 aliens following E.L.'s recipe:

```
set.seed(123)
Alien2 <- tibble(humans_eaten = rep(0:6, each = 5),
                 prob_departure = plogis(-2.197 + 1.5*humans_eaten),
                 departed = rbinom(n = 35, size = 1, prob = prob_departure))
```

GLM — data generating process: data check

```
Alien2 |> slice(1:17) |> print(n = Inf)
```

```
# A tibble: 17 x 3
```

	humans_eaten <int>	prob_departure <dbl>	departed <int>
1	0	0.100	0
2	0	0.100	0
3	0	0.100	0
4	0	0.100	0
5	0	0.100	1
6	1	0.332	0
7	1	0.332	0
8	1	0.332	1
9	1	0.332	0
10	1	0.332	0
11	2	0.691	0
12	2	0.691	1
13	2	0.691	1
14	2	0.691	1
15	2	0.691	1
16	3	0.909	1
17	3	0.909	1

```
Alien2 |> slice(18:35) |> print(n = Inf)
```

```
# A tibble: 18 x 3
```

	humans_eaten <int>	prob_departure <dbl>	departed <int>
1	3	0.909	1
2	3	0.909	1
3	3	0.909	0
4	4	0.978	1
5	4	0.978	1
6	4	0.978	1
7	4	0.978	0
8	4	0.978	1
9	5	0.995	1
10	5	0.995	1
11	5	0.995	1
12	5	0.995	1
13	5	0.995	1
14	6	0.999	1
15	6	0.999	1
16	6	0.999	1
17	6	0.999	1
18	6	0.999	1

GLM — model fitting: frequentist approach

In practice, model parameters are **unknown**, so we have to estimate them
(God mode is again switched off)

```
fit_alien_glm <- glm(departed ~ humans_eaten, data = Alien2, family = binomial(link = "logit"))
```

NB:

- we used `glm()` rather than `lm()`
- we have a new term `binomial(link = "logit")` to indicate that the data are generated via a binomial distribution and to indicate which link the fit must use²

²the link is the function that converts the linear predictor into the scale of the response, or as here into a probability (more on that later)

GLM — model fitting: frequentist approach

The summary table is not very different from those produced with `lm()` either

```
summary(fit_alien_glm)
```

Call:

```
glm(formula = departed ~ humans_eaten, family = binomial(link = "logit"),  
     data = Alien2)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.5952	0.8129	-1.962	0.04973 *
humans_eaten	1.0001	0.3411	2.932	0.00337 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 43.574 on 34 degrees of freedom
Residual deviance: 27.667 on 33 degrees of freedom
AIC: 31.667

Number of Fisher Scoring iterations: 5

GLM — model fitting: Bayesian approach with {brms}

Fitting GLM with {brms} is slightly different since it uses a specific family for binary data³

```
fit_alien_glm_brms <- brm(departed ~ humans_eaten, data = Alien2,  
  family = bernoulli(link = "logit"),  
  iter = 10000) # more needed than default this time  
summary(fit_alien_glm_brms) # or fixef(fit_alien_lm_brms) for a focus on coefficients
```

```
Family: bernoulli  
Links: mu = logit  
Formula: departed ~ humans_eaten  
Data: Alien2 (Number of observations: 35)  
Draws: 4 chains, each with iter = 10000; warmup = 5000; thin = 1;  
total post-warmup draws = 20000
```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	-1.81	0.87	-3.67	-0.28	1.00	11904	11479
humans_eaten	1.09	0.35	0.50	1.89	1.00	7029	8505

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

³the Bernoulli distribution is a special case of the binomial distribution when there is only 1 trial

GLM — model fitting: Bayesian approach with {nimble}

Defining the GLM by hand with {nimble} is relatively simple...

```
code_glm <- nimbleCode({  
  for (i in seq_len(N)) {  
    # model mimicking the data generating process  
    logit(p[i]) <- intercept + slope*humans_eaten[i]  
    departed[i] ~ dbin(p[i], 1)  
  }  
  # prior  
  intercept ~ dnorm(mean = 0, sd = 2)  
  slope ~ dnorm(mean = 0, sd = 2)  
})  
  
constants_glm <- list(N = nrow(Alien2),  
                      humans_eaten = Alien2$humans_eaten)  
  
inits_glm <- list(intercept = 0, slope = 0)
```

GLM — model fitting: Bayesian approach (prior predictive checking)

... but again we need to make sure the priors are sensible (here it would be even more difficult to foresee)

```
set.seed(123)
n_sims2 <- 100


prior_intercept2 <- rnorm(n = n_sims2, mean = 0, sd = 2)
prior_slope2 <- rnorm(n = n_sims2, mean = 0, sd = 2)

expand_grid(sim = seq_len(n_sims2),
            humans_eaten = seq(0, 6, by = 0.1)) |>
  mutate(prob = plogis(prior_intercept2[sim] +
                       prior_slope2[sim] * humans_eaten)) -> prior_pred2

ggplot(prior_pred2) +
  geom_histogram(aes(x = prob))

ggplot(prior_pred2) +
  geom_line(aes(y = prob, x = humans_eaten, group = sim))
```

- it looks OK



GLM — model fitting: Bayesian approach with {nimble}

We can now fit the model with {nimble}

```
fit_alien_glm_nimb <- nimbleMCMC(code = code_glm,  
                                constants = constants_glm,  
                                data = list(departed = Alien2$departed),  
                                inits = inits_glm, nchains = 4, setSeed = 1:4,  
                                nburnin = 4000)  
fit_alien_glm_nimb_d <- as_draws(fit_alien_glm_nimb)  
summary(fit_alien_glm_nimb_d)
```

```
fit_alien_glm_nimb_d <- as_draws(fit_alien_glm_nimb)  
summary(fit_alien_glm_nimb_d)
```

A tibble: 2 x 10

	variable	mean	median	sd	mad	q5	q95	rhat	ess_bulk	ess_tail
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	intercept	-1.46	-1.42	0.738	0.724	-2.73	-0.318	1.00	1326.	2476.
2	slope	1.00	0.977	0.321	0.312	0.526	1.57	1.00	1256.	2159.

GLM — model fitting results: estimates compared

Let's compare the results

`glm()`

```
# A tibble: 2 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)   -1.60        0.813    -3.43    -0.143
2 humans_eaten    1.00        0.341     0.440     1.82
```

`brms::brm()`

```
# A tibble: 2 x 5
  Parameter      Estimate Est.Error   Q2.5   Q97.5
  <chr>          <dbl>     <dbl>  <dbl>  <dbl>
1 Intercept     -1.81      0.866  -3.67  -0.281
2 humans_eaten    1.09      0.354   0.499   1.89
```

`nimble::nimbleMCMC()`

```
# A tibble: 2 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>      <dbl> <dbl>  <dbl>  <dbl>
1 intercept -1.46  0.738  -3.06  -0.118
2 slope      1.00  0.321   0.440   1.71
```

GLM — model fitting results: estimates compared

Let's compare the results

`glm()`

```
# A tibble: 2 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)   -1.60        0.813    -3.43    -0.143
2 humans_eaten    1.00        0.341     0.440     1.82
```

`brms::brm()`

```
# A tibble: 2 x 5
  Parameter      Estimate Est.Error   Q2.5   Q97.5
  <chr>          <dbl>     <dbl>  <dbl>  <dbl>
1 Intercept    -1.81      0.866  -3.67  -0.281
2 humans_eaten  1.09      0.354   0.499   1.89
```

`nimble::nimbleMCMC()`

```
# A tibble: 2 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>      <dbl> <dbl>  <dbl>  <dbl>
1 intercept -1.46  0.738  -3.06  -0.118
2 slope      1.00  0.321   0.440   1.71
```

- the estimates are in agreement and all intervals include ground truth values

GLM — model fitting results: estimates compared

Let's compare the results

`glm()`

```
# A tibble: 2 x 5
  Parameter      Estimate `Std. Error` `2.5 %` `97.5 %`
  <chr>          <dbl>      <dbl>    <dbl>    <dbl>
1 (Intercept)    -1.60        0.813    -3.43    -0.143
2 humans_eaten     1.00        0.341     0.440     1.82
```

`brms::brm()`

```
# A tibble: 2 x 5
  Parameter      Estimate Est.Error   Q2.5   Q97.5
  <chr>          <dbl>     <dbl>  <dbl>  <dbl>
1 Intercept     -1.81      0.866  -3.67  -0.281
2 humans_eaten   1.09      0.354   0.499   1.89
```

`nimble::nimbleMCMC()`

```
# A tibble: 2 x 5
  variable    mean    sd `2.5%` `97.5%`
  <chr>      <dbl> <dbl>  <dbl>  <dbl>
1 intercept -1.46  0.738  -3.06  -0.118
2 slope      1.00  0.321   0.440   1.71
```

- the estimates are in agreement and all intervals include ground truth values
- we never obtained -2.197 (intercept) or 1.5 (slope) because, once again, a sample is unlikely to reflect perfectly the population-level values

GLM — model fitting: interlude on link functions

What is the so-called *logit* (or logistic unit, or log-odds)?

```
fit_alien_glm <- glm(departed ~ humans_eaten, data = Alien2, family = binomial(link = "logit"))
```

- it is the inverse of the standard logistic function, and is defined as: $\text{logit}(p) = \log \frac{p}{1-p}$

proof:

$$\frac{1}{1+e^{-\text{logit}(p)}} = \frac{1}{1+e^{-\log \frac{p}{1-p}}} = \frac{1}{1+\frac{1}{e^{\log \frac{p}{1-p}}}} = \frac{1}{1+\frac{1}{\frac{p}{1-p}}} = \frac{1}{1+\frac{1-p}{p}} = \frac{1}{\frac{p+1-p}{p}} = \frac{1}{\frac{1}{p}} = p$$

GLM — model fitting: interlude on link functions

What is the so-called *logit* (or logistic unit, or log-odds)?

```
fit_alien_glm <- glm(departed ~ humans_eaten, data = Alien2, family = binomial(link = "logit"))
```

- it is the inverse of the standard logistic function, and is defined as: $\text{logit}(p) = \log \frac{p}{1-p}$


proof:

$$\frac{1}{1+e^{-\text{logit}(p)}} = \frac{1}{1+e^{-\log \frac{p}{1-p}}} = \frac{1}{1+\frac{1}{e^{\log \frac{p}{1-p}}}} = \frac{1}{1+\frac{1}{\frac{p}{1-p}}} = \frac{1}{1+\frac{1-p}{p}} = \frac{1}{\frac{p+1-p}{p}} = \frac{1}{\frac{1}{p}} = p$$

- it is the scale in which the linear predictor of the current GLM is expressed since the logit is allowed to vary between $-\infty$ and ∞ while p is constrained to vary between 0 and 1
- other link functions are possible for the binomial family (e.g. cauchit, probit), they are S-shaped too, but not identical


GLM — model fitting: interlude on link functions

Some link functions used for the binomial family



GLM — model fitting: interlude on link functions

Some link functions used for the binomial family



NB:

- the logit is a so-called *canonical link function*, the one with the best mathematical properties among alternatives, which is why it is the default of `binomial()`
- in Bayesian statistics there are fewer benefits to stick to canonical links compared to frequentist statistics, but the MCMC sampling can be difficult for some non-canonical links and how to interpret parameters can be even less straightforward than what we are about to see in a few slides → if you begin, stick to logit for binomial (& log for Poisson GLM)

GLM — Poisson GLM, data

Another very useful GLM is the Poisson GLM, which is used to fit count data

This time, we will directly use real data, the ones about the birds in Switzerland which you should already have downloaded:

```
territories <- read_csv("data_MHB/data-territories.csv")
species_count <- read_csv("data_MHB/header-species-count.csv")

territories |>
  summarise(territories_sum = sum(territories_total),
            .by = c(nameeg, namelt)) |>
  slice_max(territories_sum, n = 6) |>
  pull(namelt) -> top6sp

territories |>
  filter(namelt %in% top6sp & year == 2000) |>
  summarise(territories_top6 = sum(territories_total),
            .by = c(nameeg, samplearea_id)) |>
  left_join(species_count |> filter(year == 2000)) -> territories_top6_2000
```

GLM — Poisson GLM, data

Another very useful GLM is the Poisson GLM, which is used to fit count data

This time, we will directly use real data, the ones about the birds in Switzerland which you should already have downloaded:

```
territories_top6_2000
```

```
# A tibble: 1,584 x 12
  nameeg samplearea_id territories_top6 `x-koord` `y-koord` lon lat year route_length
  <chr>      <dbl>          <dbl>      <dbl>      <dbl> <dbl> <dbl> <dbl>      <dbl>
1 Commo~      268            0  2725500  1182500  9.08  46.8  2000      4444
2 Commo~      269           42  2497500  1118500  6.11  46.2  2000      6232
3 Commo~      270            0  2725500  1214500  9.09  47.1  2000      4081
4 Commo~      271            8  2503500  1134500  6.18  46.4  2000      4483
5 Commo~      272            0  2725500  1230500  9.10  47.2  2000      4939
# i 1,579 more rows
# i 3 more variables: median_elevation <dbl>, n_visits <dbl>, n_recorded_species <dbl>
```

GLM — Poisson GLM, fit

Let's model the number of recorded number of territories occupied at a sample area (in 2000) according to the median elevation and the number of visits performed:

```
fit_bird_glm <- glm(territories_top6 ~ median_elevation + n_visits,  
  data = territories_top6_2000,  
  family = poisson(link = "log"))  
summary(fit_bird_glm)
```

Call:

```
glm(formula = territories_top6 ~ median_elevation + n_visits,  
    family = poisson(link = "log"), data = territories_top6_2000)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.58495868	0.18387349	-19.50	<2e-16 ***
median_elevation	-0.00018697	0.00001731	-10.80	<2e-16 ***
n_visits	2.05716030	0.05951868	34.56	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 21010 on 1583 degrees of freedom
Residual deviance: 17144 on 1581 degrees of freedom
AIC: 21733


GLM — predictions and interpretation

At the scale of the linear predictor (here logit and then log), the interpretation is exactly as in LM

```
coef(fit_alien_glm)
```

```
(Intercept) humans_eaten  
-1.595208    1.000088
```

```
estimate_link(fit_alien_glm) |>  
plot(numeric_as_discrete = FALSE)
```



Interpretation

- on average, *the logit of departure* is -1.6 when no human is eaten
- on average, every additional human eaten increases *the logit of departure* by 1


GLM — predictions and interpretation

At the scale of the linear predictor (here logit and then log), the interpretation is exactly as in LM

```
coef(fit_bird_glm)
```

(Intercept)	median_elevation	n_visits
-3.584958680	-0.000186971	2.057160296

```
estimate_link(fit_bird_glm) |>  
plot()
```



Interpretation

- on average, *the log number of territories* is -3.58 when the median elevation is 0m and the number of visit is 0
- on average, every additional meter in median elevation increases *the log number of territories* by -0.000187
- on average, every additional visit increases *the log number of territories* by 2.06


GLM — predictions and interpretation

At the scale of the response variable, predictions and interpretation differ

```
coef(fit_alien_glm)
```

```
(Intercept) humans_eaten  
-1.595208    1.000088
```

```
estimate_relation(fit_alien_glm,  
  by = list(humans_eaten = seq(0, 6, length = 100))) |>  
plot(show_data = TRUE, numeric_as_discrete = FALSE)
```



Interpretation

- on average, every additional human eaten increases *the odds of departure* $e^{\text{slope}} = 2.72$ fold


GLM — predictions and interpretation

At the scale of the response variable, predictions and interpretation differ

```
coef(fit_bird_glm)
```

(Intercept)	median_elevation	n_visits
-3.584958680	-0.000186971	2.057160296

```
estimate_relation(fit_bird_glm) |>  
plot(show_data = FALSE) # no data points for clarity
```



Interpretation

- on average, every additional meter in median elevation increases the number of territories by $100 \times (e^{\text{slope}} - 1) = -0.0187\%$
- on average, every additional 1000 meter in median elevation increases the number of territories by $100 \times (e^{\text{slope} \times 1000} - 1) = -17.1\%$

Tests

Assumptions

Model construction

A simple habitat occupancy model

Questions?

Developer corner

Ignore this slide — (technical info for Alex)

```
[1] "1 min of R running time"
```

```
R version 4.6.1 (2026-06-24)
```

```
Platform: x86_64-redhat-linux-gnu
```

```
Running under: Fedora Linux 44 (KDE Plasma Desktop Edition)
```

```
Matrix products: default
```

```
BLAS/LAPACK: FlexiBLAS OPENBLAS-OPENMP; LAPACK version 3.12.1
```

```
attached base packages:
```

```
[1] stats      graphics  grDevices datasets  utils      methods    base
```

```
other attached packages:
```

```
[1] modelbased_0.16.0 bayesplot_1.15.0 posterior_1.7.0 nimble_1.4.2      brms_2.23.0      Rcpp_1.1.2      lubridate_1.9.5
[8] forcats_1.0.1      stringr_1.6.0      dplyr_1.2.1      purrr_1.2.2      readr_2.2.0      tidyr_1.3.2      tibble_3.3.1
[15] ggplot2_4.0.3      tidyverse_2.0.0
```

```
loaded via a namespace (and not attached):
```

```
[1] tidyselect_1.2.1      farver_2.1.2          loo_2.10.1            S7_0.2.2              fastmap_1.2.0
[6] pracma_2.4.6          TH.data_1.1-5         tensorA_0.36.2.1      bayestestR_0.18.1     digest_0.6.39
[11] estimability_2.0.0    timechange_0.4.0      lifecycle_1.0.5       StanHeaders_2.32.10   survival_3.8-9
[16] magrittr_2.0.5        compiler_4.6.1        rlang_1.3.0           tools_4.6.1           utf8_1.2.6
[21] igraph_2.3.3          yaml_2.3.12           data.table_1.18.4     knitr_1.51            labeling_0.4.3
[26] bridgesampling_1.2-1 bit_4.6.0             pkgbuild_1.4.8        curl_7.1.0            plyr_1.8.9
[31] RColorBrewer_1.1-3    multcomp_1.4-31       abind_1.4-8           numDeriv_2016.8-1.1   withr_3.0.3
[36] datawizard_1.3.1      grid_4.6.1            stats4_4.6.1          xtable_1.8-8          inline_0.3.21
[41] emmeans_2.0.4         scales_1.4.0          MASS_7.3-66           insight_1.5.2         tinytex_0.60
[46] cli_3.6.6            mvtnorm_1.4-2         crayon_1.5.3          rmarkdown_2.31        generics_0.1.4
[51] otel_0.2.0            RcppParallel_6.2.0    rstudioapi_0.19.0     reshape2_1.4.5        tzdb_0.5.0
[56] parameters_0.29.2     rstan_2.32.7          vctrs_0.7.3           parallel_4.6.1        CoprManager_0.5.8
[61] marginaleffects_0.32.0 matrixStats_1.5.0     vtables_0.7.3         V8_8.2.0              Matrix_1.7-6
[66] sandwich_3.1-3        jsonlite_2.0.0        hms_1.1.4             bit64_4.8.2           see_0.14.1
[71] glue_1.8.1            codetools_0.2-20      distributional_0.8.1  stringi_1.8.9         gtable_0.3.6
[76] QuickJSR_1.10.0       pillar_1.11.1         htmltools_0.5.9       Brodningnag_1.2-9     R6_2.6.1
[81] vroom_1.7.1           evaluate_1.0.5        lattice_0.23-1        backports_1.5.1      rstantools_2.7.0
[86] coda_0.19-4.1         gridExtra_2.3.1      nlme_3.1-170          checkmate_2.3.4       xfun_0.60
[91] zoo_1.8-9            rstanfit_2.0.2
```